

ABHISHEK YADAV

Full Stack Developer & AI Engineer | RAG Knowledge Base v2.0

INFO: This document is structured for Retrieval-Augmented Generation (RAG). Each [CHUNK: ID] section is a self-contained knowledge unit. Ingest each section as a separate embedding. Do NOT merge chunks across sections.

CHUNK 01

Personal Identity & Contact Information

[CHUNK: IDENTITY_001]

Abhishek Yadav is a passionate Full Stack Developer and AI Engineer based in Noida/Bangalore, India. He specializes in building intelligent, scalable web applications using the MERN stack combined with AI and LLM integration. He is currently seeking full-time, remote, or hybrid opportunities.

Full Name	Abhishek Yadav
Professional Title	Full Stack Developer & AI Engineer
Experience	1+ Year professional experience
Location	Noida / Bangalore, Karnataka, India
Work Preference	Remote Hybrid On-site (Bangalore / Noida)
Email	yadabhishek62@gmail.com
GitHub	github.com/yadavabhi12
LinkedIn	linkedin.com/in/abhishek-yadav-dev/

Retrieve this chunk for: Who is Abhishek? | Contact details | Location | Email

CHUNK 02

Professional Summary

[CHUNK: SUMMARY_002]

Abhishek Yadav is a results-driven Full Stack Developer and AI Engineer with a strong foundation in the MERN stack (MongoDB, Express.js, React.js, Node.js). His core expertise lies in building modern, scalable, user-centric web applications with a deep focus on integrating Artificial Intelligence, real-time systems, and automation workflows.

Production-level expertise includes:

- * Building and deploying AI-powered chatbots using RAG pipelines (Pinecone + Gemini Embeddings)
- * Engineering real-time systems with WebSocket technology for live communication
- * Developing LLM-orchestrated pipelines using LangChain-style architecture (Gemini & OpenAI)
- * Creating full-stack productivity platforms with analytics, authentication, and collaboration
- * Integrating third-party APIs: Tavily (live search), Stripe (payments), TTS, Speech Recognition

He is actively seeking roles in AI Engineering, Full Stack Development, or positions combining both. Motivated by creating intelligent digital products that solve real-world problems at scale.

Retrieve this chunk for: Tell me about Abhishek | What does he specialize in | Why hire him

CHUNK 03

Education & Academic Background

[CHUNK: EDUCATION_003]

B.Tech — Computer Science Engineering

Institute of Engineering & Technology (IET), Sitapur | AKTU | 2020 – 2024 | CGPA: 7.5 / 10

Senior Secondary — 12th Grade

Victory Inter College | UP Board | 2019 | 76%

Secondary — 10th Grade

Chaudhary Kalpu Yadav Inter College | UP Board | 2017 | 79%

Retrieve for: Education | College | Degree | CGPA | B.Tech

CHUNK 04

Work Experience — Frontend Developer Intern, Mala Foundation

[CHUNK: EXP_MALA_004]

Frontend Developer Intern

Mala Foundation | November 2024 – April 2025 | 6 Months | Remote

Abhishek contributed as a Frontend Developer at Mala Foundation, building responsive, performant, and scalable user interfaces for live web products. He focused on usability improvements and visual consistency across all device platforms.

Key Achievements:

- * Improved frontend performance and load time by 30% via React optimization, lazy loading, and code splitting
- * Built reusable, mobile-first UI components with React.js and Tailwind CSS
- * Integrated RESTful APIs and managed global app state with Redux
- * Collaborated with backend teams for seamless API communication and feature delivery
- * Worked with MongoDB and MySQL for data handling and API integration

Tech Stack:

	JavaScript		Tailwind					
React.js	ES6+	Redux	CSS	Node.js	Express.js	MongoDB	MySQL	Git

Retrieve for: Work experience | Internship | Mala Foundation | Professional history

CHUNK 05

Work Experience — Web Development Training, Internshala

[CHUNK: EXP_INTERNSHALA_005]

Web Development Training Intern

Internshala | Oct 2022 – Dec 2022 | 3 Months | Remote

Completed an intensive 3-month training in Front-End Web Development. Built real-world projects using clean, maintainable code and modern responsive design techniques.

Key Learnings:

- * Developed responsive, interactive websites with HTML5, CSS3, Bootstrap, and JavaScript
- * Built a personal portfolio and product landing page as practical deliverables
- * Gained proficiency in Git version control and efficient coding workflows

Tech Stack:

HTML	CSS	JavaScript	Bootstrap	Tailwind CSS	Git
------	-----	------------	-----------	--------------	-----

CHUNK 06

Project — Krishana AI: Real-Time Context-Aware Chat System

[CHUNK: PROJECT_KRISHANA_006]

Krishana AI — Real-Time Human-Like Conversational AI System

AI Engineering Project | March 2026 – Present | LIVE & DEPLOYED

Krishana AI is Abhishek's flagship AI project — an advanced real-time conversational system deployed on his portfolio website so recruiters and visitors can chat directly with an AI agent that answers questions about Abhishek's skills, experience, and projects. It uses contextual memory, RAG-based retrieval, and live data integration to deliver accurate, human-like responses.

Core Architecture & Features:

- * RAG Pipeline: Google Gemini Embeddings + Pinecone vector DB for portfolio-specific, high-accuracy Q&A;
- * LangChain-style orchestration combining LLM (Gemini) + retrieval + real-time search for intelligent answers
- * Context-aware memory system for natural multi-turn conversations with continuity
- * Real-time data fetching via Tavily API for live queries (trends, news, weather)
- * Intelligent tool-selection: system decides when to use LLM alone, RAG retrieval, or live Tavily search
- * Multilingual: responds in Hindi, English, or Hinglish matching user's language
- * Human-like tone adaptation with conversational flow and emotional intelligence
- * WebSocket real-time message delivery for instant streaming responses

Technical Workflow:

User Message -> Language Detection -> Tool Selection (LLM / RAG / Tavily) -> Context Retrieval from Pinecone -> Prompt Construction (Gemini) -> Response Generation -> WebSocket Delivery -> Frontend Render

Links:

GitHub: <https://github.com/yadavabhi12/livechat-ai>

Live Demo: <https://livechat-ai.netlify.app/>

Tech Stack:

					Gemini		Tavily		
React.js	Node.js	Express.js	WebSocket	Pinecone	LLM	LangChain	API	MongoDB	Redux

Retrieve for: AI project | Chatbot | RAG | Krishana AI | Portfolio bot | LLM project

CHUNK 07

Project — ARIA: Multimodal Voice AI Assistant

[CHUNK: PROJECT_ARIA_007]

ARIA Voice AI — Multimodal AI Assistant

AI Engineering Project | 2025 – Present | Active Development

ARIA is a real-time voice and chat AI assistant with a Jarvis-like UI inspired by Iron Man. It is Abhishek's most ambitious ongoing project — combining speech recognition, LLM reasoning, real-time data fetching, text-to-speech output, and multilingual support into a single system.

Core Features:

- * Speech Recognition: Frontend captures voice input, converts speech to text via Speech Recognition API, sends to backend
- * LLM Processing: Gemini LLM handles general questions and natural conversation
- * Tool-Calling: Intelligently routes — LLM for normal questions, Tavily API for live web search
- * TTS Output: AI responses converted to natural speech with male/female voice options

- * Multilingual: Hindi input -> Hindi response | English -> English | Hinglish -> auto-detects
- * Voice + Text dual mode: users interact via speech or by typing
- * Jarvis-Like UI: Iron Man inspired futuristic animated interface
- * Future plans: computer automation tasks, calendar management, file operations via tool-calling

Tech Stack:

React.js	Node.js	Express.js	Gemini LLM	Mistral AI	Tavily API	Speech Recognition API	WebSocket	TTS API
----------	---------	------------	------------	------------	------------	------------------------	-----------	---------

Links:

GitHub: https://github.com/yadavabhi12/AI_Voice_Agent

Retrieve for: Voice AI | ARIA | Voice assistant | Speech recognition project | Jarvis-like AI

CHUNK 08

Project — Focus Engine: Full-Stack Productivity Platform

[CHUNK: PROJECT_FOCUS_008]

Focus Engine — Smart Productivity & Collaboration Platform

Full-Stack Project | April 2026 – Present | LIVE & DEPLOYED

Focus Engine is a comprehensive full-stack productivity platform designed to solve real workflow challenges for students and professionals. It combines task management, focus tracking, real-time collaboration, and data-driven analytics into one unified system.

Core Features:

- * Secure multi-user authentication with JWT and isolated user data
- * Full task management: CRUD with reminders, priority levels, status tracking
- * Real-time WebSocket chat: private messages, group chat, file sharing, live notifications
- * Pomodoro-based focus engine with task linking, session tracking, distraction-free workflow
- * Analytics dashboard: visual charts, productivity trends (daily/weekly/monthly), exportable reports (PDF, JSON, CSV)
- * LinkedIn-style profile system with achievements, statistics, and personalization
- * Admin controls with role-based permissions for group management
- * Streak-based habit tracking for long-term consistency

Links:

GitHub: <https://github.com/yadavabhi12/abhishekyadav-focus-engine>

Live Demo: <https://productivityfocus.netlify.app/>

Tech Stack:

React.js	Redux	Node.js	Express.js	WebSocket	MongoDB	MySQL	Tailwind CSS	Chart.js
----------	-------	---------	------------	-----------	---------	-------	--------------	----------

Retrieve for: Productivity app | Focus Engine | Task manager | Real-time chat project

CHUNK 09

Project — VoiceCraft Pro: AI Text-to-Speech App

[CHUNK: PROJECT_TTS_009]

VoiceCraft Pro — AI Text-to-Speech with MP3 Download

AI Tool Project | 2025 | LIVE

VoiceCraft Pro transforms text into lifelike Hindi and English speech with real-time audio generation and instant MP3 downloads. Demonstrates ability to integrate AI APIs into clean, user-friendly interfaces.

Features:

- * Converts text to natural-sounding speech in Hindi and English
- * Real-time audio generation with instant browser playback
- * One-click MP3 download of generated audio
- * Clean responsive UI built with React.js

Links:

GitHub: https://github.com/yadavabhi12/mern_text_to_speech

Live Demo: <https://merntexttospeech.netlify.app/>

React.js	Node.js	Express.js	TTS API
----------	---------	------------	---------

Project — FlowMind AI: Visual Prompt Builder

[CHUNK: PROJECT_FLOWMIND_010]

FlowMind AI — Node-Based Visual AI Workflow Builder

AI Tool Project | 2025 | LIVE

FlowMind AI is a node-based visual workflow builder where users connect prompt nodes to AI models, execute real-time LLM responses, and store results. Inspired by n8n and LangGraph, it demonstrates Abhishek's deep understanding of agentic AI pipeline design and visual programming.

- * Drag-and-drop node interface using React Flow for building visual prompt pipelines
- * Real-time execution of AI workflows powered by Gemini API
- * Results storage and retrieval with MongoDB
- * Modular architecture supporting addition of new node types

React.js	React Flow	Node.js	Express.js	Gemini API	MongoDB
----------	------------	---------	------------	------------	---------

CHUNK 11

Project — IntelData Datastore: Data Management Platform

[CHUNK: PROJECT_INTELDATA_011]

IntelData Datastore — Multi-User Structured Data Management Platform

Full-Stack Project | 2024 | LIVE

IntelData is a secure full-stack platform acting as a 'private Google Drive for structured data'. It uses a 3-tier architecture with Java Spring Boot backend, demonstrating Abhishek's versatility beyond Node.js.

- * 3-Tier Architecture: React.js frontend, Java Spring Boot backend, MongoDB database
- * JWT-based auth with role-based access control (Admin/User roles)
- * Full CRUD for datasets with backend validation
- * Real-time data tables with filtering, pagination, and summary statistics
- * 30% improvement in page load time through React optimization and backend caching

React.js	Tailwind CSS	Java Spring Boot	Spring Security	MongoDB	RESTful APIs
----------	--------------	------------------	-----------------	---------	--------------

Retrieve for: IntelData | Spring Boot project | Java backend | Data management app

CHUNK 12

Project — FoodDash Pro: Food Delivery Platform

[CHUNK: PROJECT_FOOD_012]

FoodDash Pro — Dual Panel Food Delivery System

Full-Stack MERN Project | 2024 | LIVE

FoodDash Pro is a complete e-commerce food delivery platform with separate user and admin panels. It covers the full customer journey from browsing to order placement, payment, and tracking.

- * Responsive UI: menu browsing, cart management, real-time order tracking
- * Admin panel: item management, order processing, status updates
- * Stripe payment gateway integration for secure online payments
- * JWT authentication for users and admin

Links:

GitHub: <https://github.com/yadavabhi12/food-delevery>

React.js	Redux	Node.js	Express.js	MongoDB	Stripe API	Tailwind CSS	JWT
----------	-------	---------	------------	---------	------------	--------------	-----

Technical Skills & Expertise

[CHUNK: SKILLS_013]

Complete technical skills reference for Abhishek Yadav, organized by domain. Use this chunk for any question about his technical capabilities or specific technologies.

Frontend

React.js (Advanced)	JavaScript ES6+ (Advanced)	HTML5/CSS3 (Advanced)	Tailwind CSS / Bootstrap (Advanced)	Redux / Context API (Proficient)	Responsive & Accessible UI (Advanced)
------------------------	-------------------------------	--------------------------	---	-------------------------------------	---

Backend

Node.js (Advanced)	Express.js (Advanced)	RESTful API Design (Advanced)	WebSocket / Real-Time Systems (Proficient)	JWT Auth & Authorization (Proficient)
--------------------	--------------------------	----------------------------------	--	---

AI / LLM / ML

Gemini API - LLM Integration (Proficient)	OpenAI API (Proficient)	LangChain & LangGraph (Learning/Applying)	Pinecone Vector DB (Proficient)	RAG Pipeline Implementation (Proficient)	Tavily API - Live Search (Proficient)	Mistral AI (Applying)	Python, NumPy, Pandas (Learning)
--	----------------------------	--	---------------------------------------	---	--	--------------------------	--

Databases

MongoDB (Proficient)	MySQL (Proficient)	Spring Boot + Spring Security (Familiar)
----------------------	--------------------	---

Tools & Cloud

Git & GitHub (Advanced)	Postman (Advanced)	VS Code (Advanced)	AWS Cloud (Learning)	n8n Automation (Learning)
----------------------------	--------------------	--------------------	-------------------------	------------------------------

Problem Solving

LeetCode: 200+ Problems Solved	Java - DSA & OOP (Proficient)
--------------------------------	-------------------------------

Retrieve for: Skills | Technologies | Does he know X | React experience | AI skills

CHUNK 14

All Projects — Consolidated Reference Table

[CHUNK: PROJECTS_ALL_014]

Abhishek has completed 7+ shipped projects spanning AI systems, full-stack web apps, productivity tools, and API integrations. All major projects are live and deployed.

Project	Type	Status	Live URL	GitHub
Krishana AI	AI Chat System	Live	livechat-ai.netlify.app	github.com/yadavabhi12/livechat-ai
Focus Engine	Productivity Platform	Live	productivityfocus.netlify.app	github.com/yadavabhi12/abhishek-yadav-focus-engine
ARIA Voice AI	Voice AI Assistant	In Dev	—	github.com/yadavabhi12/AI_Voice_Agent
VoiceCraft Pro	TTS App	Live	merntexttospeech.netlify.app	github.com/yadavabhi12/mern_text_to_speech
FlowMind AI	Visual Prompt Builder	Live	—	—
IntelData	Data Mgmt Platform	Live	—	—
FoodDash Pro	Food Delivery App	Live	—	github.com/yadavabhi12/food-delivery

Retrieve for: All projects | Project list | How many projects | Deployed work | Project links

CHUNK 15

Career Goals, Learning Roadmap & Future Plans

[CHUNK: GOALS_015]

Currently Learning (2025–2026):

- * AI Engineering: LangChain, LangGraph, autonomous AI agents with vector search
- * Machine Learning: NumPy, Pandas, Scikit-learn for data analysis and ML automation
- * Cloud & DevOps: AWS services — targeting AWS certification
- * System Design: Microservices, design patterns, high-performance architecture
- * n8n Automation: Building AI-powered automation workflows
- * DSA: Targeting 500+ LeetCode problems

Short-Term Goals (2025–2026):

- * Earn AWS Cloud Practitioner / Solutions Architect certification
- * Build and deploy a production SaaS product with integrated AI agents
- * Complete ARIA with full tool-calling and computer automation capabilities
- * Reach 500+ LeetCode problems solved

Long-Term Career Aspirations:

- * AI Product Engineer: Build and deploy intelligent, AI-first products solving real-world problems
- * System Design Specialist: Architect large-scale, fault-tolerant backend systems
- * Entrepreneur: Create a generative AI platform for personalized task automation
- * Open Source & Mentorship: Contribute to AI open-source and mentor developers

Retrieve for: Goals | Future plans | What is Abhishek learning | AWS | Career aspirations

CHUNK 16

Key Professional Achievements

[CHUNK: ACHIEVEMENTS_016]

- * Deployed Krishana AI — end-to-end AI chat system with RAG pipeline (Pinecone + Gemini), LLM orchestration, WebSocket, multilingual support — live on portfolio
- * Built ARIA Voice AI: multimodal assistant with Speech-to-Text, LLM reasoning, tool-calling (Tavily), TTS, Hindi/English/Hinglish support
- * Improved frontend performance by 30% at Mala Foundation via lazy loading, code splitting, and React optimization
- * IntelData Datastore: 30% page load improvement through React optimization and backend caching
- * Built Focus Engine — production-level multi-user productivity platform with real-time chat and analytics — live deployed
- * Completed 200+ LeetCode problems demonstrating strong DSA skills
- * Shipped 7+ full-stack and AI projects — all deployed on Netlify and publicly accessible
- * Consistently implemented JWT authentication and data privacy mechanisms across all projects

Retrieve for: Achievements | Accomplishments | Why hire | What makes him stand out | Results

CHUNK 17

Personal Interests & Personality

[CHUNK: PERSONAL_017]

Outside of professional work, Abhishek maintains a balanced lifestyle that fosters teamwork, creativity, and continuous learning.

- * Sports: Badminton and Volleyball — reflects teamwork and competitive mindset
- * Entertainment: Movies for relaxation and creative inspiration
- * Continuous Learning: Follows latest AI, automation, and web tech trends via blogs and research
- * Side Projects: Treats personal builds as real-world learning experiments

Personality: Detail-oriented, self-motivated, collaborative, passionate about problem-solving, and driven by building meaningful digital products.

Retrieve for: Hobbies | Personal info | What does he do outside work | Personality

Frequently Asked Questions — Recruiter Q&A;

[CHUNK: FAQ_018]

Pre-formatted Q&A; pairs for common recruiter and visitor questions. The chatbot should prioritize this chunk for direct question matching.

Q: Is Abhishek available for work?

A: Yes. Abhishek is actively seeking full-time, remote, or hybrid opportunities in Full Stack Development or AI Engineering. He is also open to freelance projects. Contact: yadabhishek62@gmail.com or linkedin.com/in/abhishek-yadav-dev/

Q: What is Abhishek's most impressive project?

A: Krishana AI is his most technically complex — a full-stack real-time AI system with RAG pipeline (Pinecone + Gemini Embeddings), LangChain-style orchestration, WebSocket delivery, multilingual support, and live Tavily data integration. Deployed live at livechat-ai.netlify.app.

Q: Does Abhishek have real AI experience?

A: Yes, hands-on production experience with: Gemini API, OpenAI API, LangChain-style pipelines, Pinecone vector databases, full RAG implementation, Tavily real-time search, Mistral AI, Speech Recognition, and TTS integration. He is actively building autonomous AI agents with LangGraph.

Q: What is the RAG chatbot on his portfolio?

A: Krishana AI is integrated into Abhishek's portfolio so recruiters can directly chat with an AI agent. The bot uses RAG — it retrieves relevant information from a Pinecone vector database using Gemini embeddings and generates accurate, context-aware answers about his skills, experience, and projects.

Q: Can Abhishek work with Java or Spring Boot?

A: Yes. Abhishek built IntelData Datastore using Java Spring Boot, Spring Security, and MongoDB as backend. He is also proficient in Java for Data Structures and Algorithms.

Q: What makes Abhishek different from other developers?

A: He uniquely combines Full Stack + AI Engineering. He builds production-level AI systems — not just basic CRUD apps. He has shipped RAG pipelines, LLM orchestration, voice AI with multilingual support, and real-time systems that are all deployed and publicly accessible. Every project demonstrates real engineering capability, not just tutorial reproductions.

Q: How can I see Abhishek's live work?

A: Focus Engine: productivityfocus.netlify.app | Krishana AI Chat: livechat-ai.netlify.app | VoiceCraft TTS: merntexttospeech.netlify.app | GitHub: github.com/yadavabhi12

Q: What is his salary expectation?

A: This is not publicly disclosed. For salary/CTC discussions, contact Abhishek directly at yadabhishek62@gmail.com or via LinkedIn: linkedin.com/in/abhishek-yadav-dev/

RAG Pipeline Ingestion & Configuration Guide

[CHUNK: RAG_GUIDE_019]

This section is a technical reference for ingesting this document into the RAG pipeline.

Chunking Strategy:

- * Split by SECTION HEADERS — each [CHUNK: ID] block is one independent embedding unit
- * Do NOT merge chunks across sections during ingestion
- * Recommended chunk overlap: 0 (sections are already self-contained)
- * Embedding model: Google text-embedding-004 or text-embedding-003

Metadata per chunk:

- * chunk_id: Use [CHUNK: ID] value — e.g., IDENTITY_001, PROJECT_KRISHANA_006
- * section_type: identity | summary | education | experience | project | skills | goals | faq
- * project_name: (for project chunks) krishana_ai | focus_engine | aria | voicecraft | etc.
- * source: abhishek_yadav_rag_v2

Retrieval Configuration:

- * Top-K: 3 to 5 chunks per query
- * Similarity threshold: minimum 0.75 cosine similarity
- * FAQ_018 chunk should have boosted retrieval weight for direct question matching
- * For 'contact' queries: always include IDENTITY_001
- * For 'projects' queries: include PROJECTS_ALL_014 + specific project chunk

Privacy & Language Notes:

- * Phone number is intentionally excluded from this document — do not include in vector DB
- * If user query is in Hindi: translate to English before embedding, respond in Hindi
- * If query is Hinglish: detect dominant language and process accordingly

Document authored by Abhishek Yadav for Krishana AI Portfolio RAG System — Version 2.0 (2026)